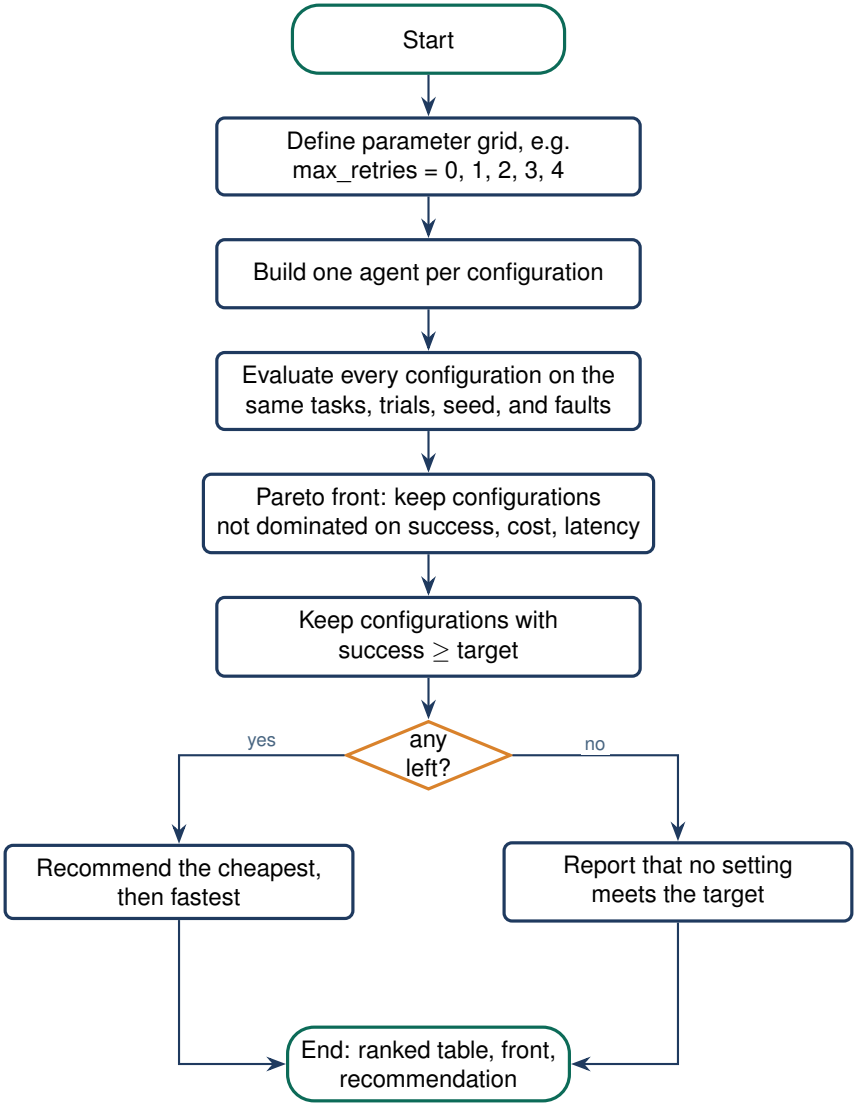
Start
Define parameter grid, e.g. max_retries = 0, 1, 2, 3, 4
Build one agent per configuration
Evaluate every configuration on the same tasks, trials, seed, and faults
Pareto front: keep configurations not dominated on success, cost, latency
Keep configurations with success ≥ target
any left?
yes
no
Recommend the cheapest, then fastest
Report that no setting meets the target
End: ranked table, front, recommendation