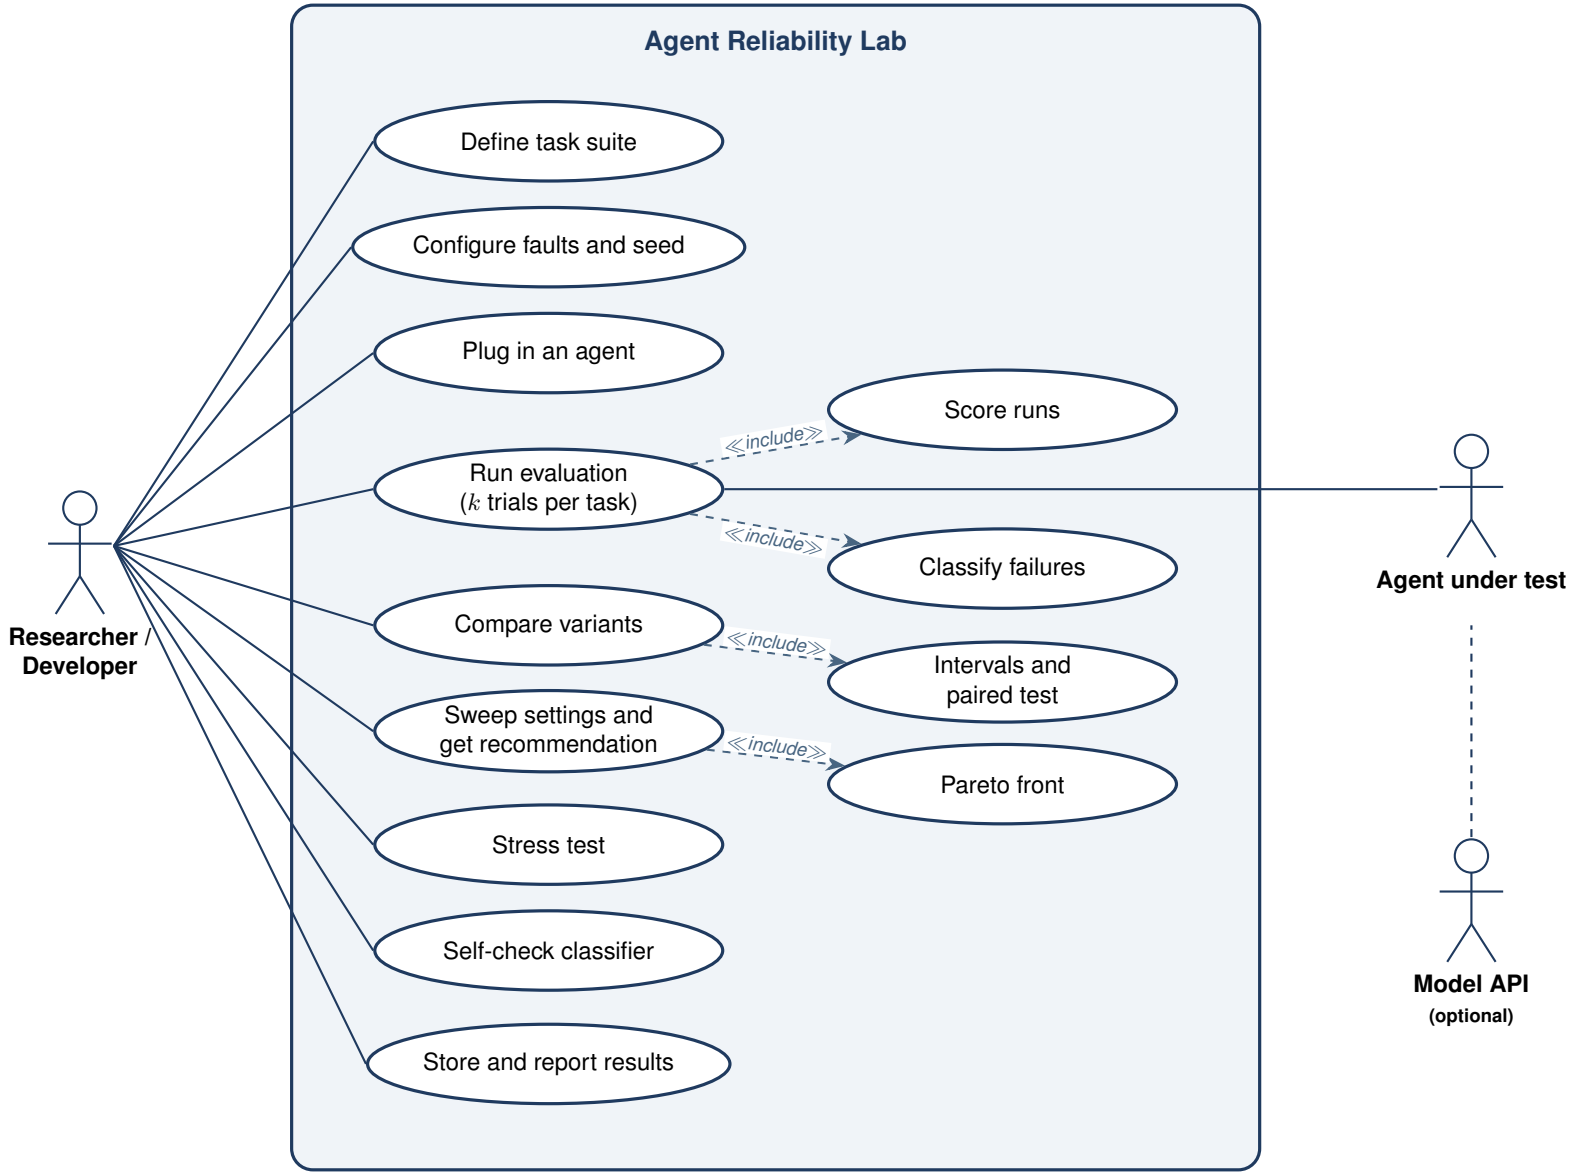
Agent Reliability Lab
Define task suite
Configure faults and seed
Plug in an agent
Score runs
Run evaluation
(k trials per task)
Classify failures
Compare variants
Intervals and
paired test
Sweep settings and
get recommendation
Pareto front
Stress test
Self-check classifier
Store and report results
«include»
«include»
«include»
«include»
Researcher /
Developer
Agent under test
Model API
(optional)